

## Q1: Import / Load

```
In [15]: # Step 1: Import libraries and set random seeds
import numpy as np
import tensorflow as tf
from tensorflow.keras import Sequential, layers
from tensorflow.keras.datasets import fashion_mnist
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

# Set random seed for reproducibility
tf.random.set_seed(42)
np.random.seed(42)

(X_train, y_train), (X_test, y_test) = fashion_mnist.load_data()

label_names = ["T-shirt/top", "Trouser", "Pullover", "Dress", "Coat",
               "Sandal", "Shirt", "Sneaker", "Bag", "Ankle boot"]

x_train = x_train / 255.0
x_test = x_test / 255.0

# Visualize a few samples
plt.figure(figsize=(10,10))
for i in range(9):
    plt.subplot(3,3,i+1)
    plt.imshow(X_train[i], cmap='gray')
    plt.title(label_names[y_train[i]])
    plt.axis('off')
plt.tight_layout()
plt.show()
```



## one hot encode

```
In [16]: # Step 1 (continued): Flatten and normalize, then one-hot encode
X_train = X_train.reshape(-1, 28*28).astype('float32') / 255.0
X_test  = X_test.reshape(-1, 28*28).astype('float32') / 255.0

y_train = to_categorical(y_train, 10)
y_test  = to_categorical(y_test, 10)
```

## Define and compile model

```
In [17]: def create_model():
    model = Sequential([
        layers.Input(shape=(784,)),
        layers.Dense(512, activation='relu'),
        layers.Dropout(0.2),
```

```

        layers.Dense(256, activation='relu'),
        layers.Dropout(0.2),
        layers.Dense(128, activation='relu'),
        layers.Dropout(0.2),
        layers.Dense(10, activation='softmax')
    ])
    return model

model = create_model()
model.summary()

model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 512)	401,920
dropout_6 (Dropout)	(None, 512)	0
dense_11 (Dense)	(None, 256)	131,328
dropout_7 (Dropout)	(None, 256)	0
dense_12 (Dense)	(None, 128)	32,896
dropout_8 (Dropout)	(None, 128)	0
dense_13 (Dense)	(None, 10)	1,290

Total params: 567,434 (2.16 MB)

Trainable params: 567,434 (2.16 MB)


Non-trainable params: 0 (0.00 B)


## Train model


```


In [18]: # Step 4: Train the model
history = model.fit(
    X_train, y_train,
    validation_data=(X_test, y_test),
    epochs=20,
    batch_size=128,
    verbose=1
)


```


Epoch 1/20  
469/469  3s 4ms/step - accuracy: 0.7953 - loss: 0.5730 - val\_accuracy: 0.8493 - val\_loss: 0.4206


Epoch 2/20  
469/469  2s 4ms/step - accuracy: 0.8526 - loss: 0.4085 - val\_accuracy: 0.8581 - val\_loss: 0.3988


Epoch 3/20  
469/469  2s 4ms/step - accuracy: 0.8648 - loss: 0.3700 - val\_accuracy: 0.8633 - val\_loss: 0.3764


Epoch 4/20  
469/469  2s 4ms/step - accuracy: 0.8751 - loss: 0.3446 - val\_accuracy: 0.8710 - val\_loss: 0.3596


Epoch 5/20  
469/469  2s 4ms/step - accuracy: 0.8802 - loss: 0.3289 - val\_accuracy: 0.8736 - val\_loss: 0.3530


Epoch 6/20  
469/469  2s 4ms/step - accuracy: 0.8837 - loss: 0.3150 - val\_accuracy: 0.8800 - val\_loss: 0.3338


Epoch 7/20  
469/469  2s 4ms/step - accuracy: 0.8874 - loss: 0.3027 - val\_accuracy: 0.8835 - val\_loss: 0.3343


Epoch 8/20  
469/469  2s 4ms/step - accuracy: 0.8921 - loss: 0.2918 - val\_accuracy: 0.8766 - val\_loss: 0.3395


Epoch 9/20  
469/469  2s 4ms/step - accuracy: 0.8934 - loss: 0.2851 - val\_accuracy: 0.8818 - val\_loss: 0.3278


Epoch 10/20  
469/469  2s 4ms/step - accuracy: 0.8978 - loss: 0.2742 - val\_accuracy: 0.8823 - val\_loss: 0.3413


Epoch 11/20  
469/469  2s 4ms/step - accuracy: 0.8985 - loss: 0.2682 - val\_accuracy: 0.8852 - val\_loss: 0.3234


Epoch 12/20  
469/469  2s 4ms/step - accuracy: 0.9007 - loss: 0.2629 - val\_accuracy: 0.8868 - val\_loss: 0.3256


Epoch 13/20  
469/469  2s 4ms/step - accuracy: 0.9022 - loss: 0.2574 - val\_accuracy: 0.8833 - val\_loss: 0.3287


Epoch 14/20  
469/469  3s 7ms/step - accuracy: 0.9038 - loss: 0.2545 - val\_accuracy: 0.8890 - val\_loss: 0.3173

Epoch 15/20  
469/469  6s 7ms/step - accuracy: 0.9065 - loss: 0.2467 - val\_accuracy: 0.8843 - val\_loss: 0.3312

Epoch 16/20  
469/469  4s 5ms/step - accuracy: 0.9078 - loss: 0.2421 - val\_accuracy: 0.8881 - val\_loss: 0.3229

Epoch 17/20  
469/469  2s 4ms/step - accuracy: 0.9108 - loss: 0.2362 - val\_accuracy: 0.8889 - val\_loss: 0.3169

Epoch 18/20  
469/469  2s 4ms/step - accuracy: 0.9112 - loss: 0.2343 - val\_accuracy: 0.8910 - val\_loss: 0.3184

Epoch 19/20  
469/469  2s 4ms/step - accuracy: 0.9125 - loss: 0.2298 - val\_accuracy: 0.9125 - val\_loss: 0.2298

uracy: 0.8901 - val\_loss: 0.3249

Epoch 20/20

469/469 ————— 2s 4ms/step - accuracy: 0.9157 - loss: 0.2231 - val\_acc

uracy: 0.8932 - val\_loss: 0.3225


## plot

```
In [19]: # Step 5: Evaluate test accuracy
test_loss, test_acc = model.evaluate(X_test, y_test, verbose=0)
print(f"Test accuracy: {test_acc*100:.2f}%")

# Plot accuracy and loss
plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.tight_layout()
plt.show()
```

Test accuracy: 89.32%



## Q2

Q2. For a given set of images with titles and without titles, which data set will have a better image classification performance? Explain and justify your answer.

The dataset with titles (labeled images) will lead to significantly better classification performance, because supervised learning relies on labeled data to compute error, adjust

model weights, and improve predictive accuracy. Unlabeled images cannot train a classifier well without additional labeling or semi-supervised techniques.